

# DLP - Week-10 - PA

1) Given an image dataset for object detection, how would you preprocess the data and create region proposals for Fast R-CNN? Consider the following code snippet:

1 point

```
import cv2
import numpy as np
from skimage.feature import hog
```

```
image = cv2.imread('path_to_image.jpg')
resized_image = cv2.resize(image, (600, 600))
hog_features = hog(resized_image, orientations=9, pixels_per_cell=(8, 8),
                   cells_per_block=(2, 2), visualize=False, multichannel=True)
```

```
def generate_proposals(image, window_size=(128, 128), step_size=32):
    proposals = []
    for y in range(0, image.shape[0] - window_size[1], step_size):
        for x in range(0, image.shape[1] - window_size[0], step_size):
            proposals.append((x, y, x + window_size[0], y + window_size[1]))
    return proposals
```

```
region_proposals = generate_proposals(resized_image)
```

Which of the following steps is NOT a part of Fast R-CNN preprocessing?

☐ (a) Resize the input image to a fixed size.

☒ (b) Extract features using HOG.

☐ (c) Generate region proposals.

☐ (d) Classify each region proposal directly.

2) Which of the following statements correctly differentiates Fast R-CNN from YOLO?

1 point

☐ (a) Fast R-CNN is a single-stage detector, while YOLO is a two-stage detector.

☒ (b) Fast R-CNN relies on external region proposals, whereas YOLO divides the image into a grid for detection.

☐ (c) YOLO uses Selective Search for region proposals, while Fast R-CNN performs end-to-end detection in one step.

☐ (d) Both Fast R-CNN and YOLO perform detection and classification in a single stage.

3) Consider the following code snippet that outlines the steps of extracting regions of interest (RoI) for object detection:

1 point

```
import cv2
import numpy as np
from sklearn.svm import SVC
from skimage.feature import hog
```

```
image = cv2.imread('dog_cat.jpg')
ss = cv2.ximgproc.segmentation.createSelectiveSearchSegmentation()
ss.setBaseImage(image)
ss.switchToSelectiveSearchFast()
region_proposals = ss.process()
```

```
## Extract features using convnet
# Step 3: Train a classifier (SVM) on the extracted features
classifier = SVC(kernel='linear')
classifier.fit(features, labels) # labels are pre-defined for training
```

Which architecture follows this approach of using region proposals, feature extraction, and a classifier for object detection?

☐ (a) Fast R-CNN

☒ (b) R-CNN (Slow R-CNN)

☐ (c) YOLO

☐ (d) Faster RCNN

4) The following code snippet demonstrates a simplified implementation of a Region Proposal Network (RPN) used in Faster R-CNN:

1 point

```
import torch
import torch.nn as nn
import torch.nn.functional as F
```

```
# Input feature map from a backbone network (e.g., ResNet)
feature_map = torch.randn(1, 512, 50, 50) # (Batch, Channels, Height, Width)
```

```
# RPN layers
```

```
class RegionProposalNetwork(nn.Module):
    def __init__(self, in_channels, num_anchors):
        super(RegionProposalNetwork, self).__init__()
        self.conv = nn.Conv2d(in_channels, 512, kernel_size=3, stride=1, padding=1)
        self.cls_layer = nn.Conv2d(512, num_anchors * 2, kernel_size=1) # Classification
        self.reg_layer = nn.Conv2d(512, num_anchors * 4, kernel_size=1) # Regression
```

→ resize  
→ hog features  
→ creating ROIs here  
this is not present in faster RCNN

YOLO → s x s → Bounding Box (grid like structure)

Slow RCNN  
SVM linear  
↑  
[ ] [ ] [ ]  
↑  
[ ] segmenta ~

```
def forward(self, x):
    x = F.relu(self.conv(x))
    cls_logits = self.cls_layer(x) # Objectness scores
    reg_deltas = self.reg_layer(x) # Bounding box adjustments
    return cls_logits, reg_deltas

# Initialize the RPN with 9 anchors per location
rpn = RegionProposalNetwork(in_channels=512, num_anchors=9)

# Generate RPN outputs
cls_logits, reg_deltas = rpn(feature_map)
print("Classification logits shape:", cls_logits.shape) # (1, 18, 50, 50)
print("Regression deltas shape:", reg_deltas.shape) # (1, 36, 50, 50)
```

obj. / obj not  
n

classification  
regression  
coordinates  
4n

faster RCNN

Which of the following architectures incorporates this type of Region Proposal Network (RPN) for object detection?

- ☐ (a) YOLO
- ☒ (b) Faster R-CNN
- ☐ (c) HOG detector
- ☐ (d) Fast R-CNN

5) Consider the following code snippet, which generates predictions and ground truth for an object detection task:

1 point

# Predicted bounding boxes and confidence scores

```
predictions = [
    {'bbox': [50, 50, 100, 100], 'score': 0.9},
    {'bbox': [30, 30, 70, 70], 'score': 0.75},
    {'bbox': [200, 200, 250, 250], 'score': 0.6}
]
```

TP = 2  
FP = 1  
FN = 0

# Ground truth bounding boxes

```
ground_truths = [
    {'bbox': [50, 50, 100, 100]},
    {'bbox': [30, 30, 70, 70]}
]
```

$$\text{Precision} = \frac{TP}{TP + FP} = \frac{2}{2 + 1} = \frac{2}{3} = 0.67$$

$$\text{Recall} = \frac{TP}{TP + FN} = \frac{2}{2 + 0} = \frac{2}{2} = 1$$

# Function to calculate IoU (Intersection over Union)

```
def calculate_iou(box1, box2):
```

```
    x1 = max(box1[0], box2[0])
    y1 = max(box1[1], box2[1])
    x2 = min(box1[2], box2[2])
    y2 = min(box1[3], box2[3])
    intersection = max(0, x2 - x1) * max(0, y2 - y1)
    area1 = (box1[2] - box1[0]) * (box1[3] - box1[1])
    area2 = (box2[2] - box2[0]) * (box2[3] - box2[1])
    union = area1 + area2 - intersection
    return intersection / union if union > 0 else 0
```

```
threshold = 0.5
```

```
true_positives = 0
```

```
false_positives = 0
```

```
false_negatives = len(ground_truths)
```

```
for pred in predictions:
```

```
    matched = False
```

```
    for gt in ground_truths:
```

```
        iou = calculate_iou(pred['bbox'], gt['bbox'])
```

```
        if iou >= threshold:
```

```
            true_positives += 1
```

```
            false_negatives -= 1
```

```
            matched = True
```

```
            break
```

```
    if not matched:
```

```
        false_positives += 1
```

# Calculate precision and recall using the precision\_recall function.

```
precision, recall = find_precision(true_positives, false_positives, true_negative, false_negative)
```

```
print(f'Precision: {precision:.2f}')
```

```
print(f'Recall: {recall:.2f}')
```

Given the predictions and ground truths, what are the precision and recall values?

- ☒ (a) Precision: 0.67, Recall: 1.00
- ☐ (b) Precision: 1.00, Recall: 1.00
- ☐ (c) Precision: 0.67, Recall: 0.67
- ☐ (d) Precision: 1.00, Recall: 0.67

6) The following code snippet demonstrates the calculation of Mean Average Precision (mAP) for an object detection task:

1 point

```
import numpy as np

# Predicted boxes with confidence scores and ground truth boxes
predictions = [
    {'bbox': [50, 50, 100, 100], 'score': 0.9, 'class': 'cat'},
    {'bbox': [30, 30, 70, 70], 'score': 0.8, 'class': 'dog'},
    {'bbox': [200, 200, 250, 250], 'score': 0.7, 'class': 'cat'}
]

# Ground truth boxes
ground_truths = [
    {'bbox': [50, 50, 100, 100], 'class': 'cat'},
    {'bbox': [30, 30, 70, 70], 'class': 'dog'}
]

# Function to calculate IoU (Intersection over Union)
def calculate_iou(box1, box2):
    x1 = max(box1[0], box2[0])
    y1 = max(box1[1], box2[1])
    x2 = min(box1[2], box2[2])
    y2 = min(box1[3], box2[3])
    intersection = max(0, x2 - x1) * max(0, y2 - y1)
    area1 = (box1[2] - box1[0]) * (box1[3] - box1[1])
    area2 = (box2[2] - box2[0]) * (box2[3] - box2[1])
    union = area1 + area2 - intersection
    return intersection / union if union > 0 else 0

# Evaluate precision and recall
tp, fp = 0, 0
matched_gt = set()
precision_recall_curve = []
for pred in predictions:
    matched = False
    for i, gt in enumerate(ground_truths):
        if gt['class'] == pred['class'] and i not in matched_gt:
            iou = calculate_iou(pred['bbox'], gt['bbox'])
            if iou >= 0.5: # IoU threshold
                tp += 1
                matched_gt.add(i)
                matched = True
                break
    if not matched:
        fp += 1

precision, recall = precision_recall(tp, fp, tn, fn)
# Calculate mAP as the average precision
mAP = np.mean([p for p, r in precision_recall_curve])
print(f'Mean Average Precision (mAP): {mAP:.2f}')
precision_recall_curve.append((precision, recall))
```

$$\begin{aligned} tp &= 2 \\ fp &= 1 \\ tn &= 0 \\ fn &= 0 \end{aligned}$$

$$\begin{aligned} \text{precision} &= 0.67 \\ \text{recall} &= 1 \end{aligned}$$

Given the predictions and ground truths, what is the Mean Average Precision (mAP)?

- ☐ (a) 0.33
- ☐ (b) 0.50
- ☒ (c) 0.67

7) Consider the following code snippet

1 point

```
import torch
import torch.nn as nn
import torch.nn.functional as F
from torchvision.ops import roi_pool
```

```
Step 1
feature_map = # Input feature map (e.g., from a CNN backbone like ResNet)
# Region proposals as (x1, y1, x2, y2)
region_proposals = torch.tensor([
    [0, 0, 7, 7],
    [2, 2, 10, 10],
    [4, 4, 14, 14]
], dtype=torch.float32)
```

```
Step 2
# ROI Pooling
pooled_features = roi_pool(
    feature_map,
    [region_proposals],
    output_size=(7, 7) # Fixed size for each region
)
```

```
Step 3
# Fully connected layer for classification and regression
class FastRCNNHead(nn.Module):
    def __init__(self, in_features, num_classes):
        super(FastRCNNHead, self).__init__()
        self.fc1 = nn.Linear(in_features, 1024)
        self.fc2 = nn.Linear(1024, num_classes + 4) # +4 for bbox regression
    def forward(self, x):
        x = F.relu(self.fc1(x))
        return self.fc2(x)
```

Which architecture incorporates the above components for object detection?

- ☐ (a) R-CNN (Slow R-CNN)
- ☒ (b) Fast R-CNN
- ☐ (c) Faster R-CNN
- ☐ (d) YOLO

fast RCNN  $\Rightarrow$  bbox + linear svm

8) Consider the following code snippet, which calculates the F-score for an object detection task:

1 point

# Predicted labels and ground truth labels

predictions = [1, 0, 1, 1, 0, 1, 0]

ground\_truth = [1, 0, 1, 0, 0, 1, 1]

# Confusion matrix components

true\_positive = sum([1 for p, g in zip(predictions, ground\_truth) if p == 1 and g == 1])

false\_positive = sum([1 for p, g in zip(predictions, ground\_truth) if p == 1 and g == 0])

false\_negative = sum([1 for p, g in zip(predictions, ground\_truth) if p == 0 and g == 1])

# F-score calculation find\_f1\_Score()

f\_score = find\_f1\_Score(true\_positive, false\_positive, false\_negative)

print(f"F-Score: {f\_score:.2f}")

Given the predictions and ground truth, what is the calculated F-score?

☐ (a) 0.67

☐ (b) 0.71

☒ (c) 0.75

☐ (d) 0.80

$$\text{precision} = \frac{3}{4} = 0.75$$

$$\text{recall} = \frac{3}{4} = 0.75$$

$$\frac{2 \times p \times r}{p + r} = \frac{2 \times 0.75 \times 0.75}{0.75 + 0.75}$$

9) Consider the following code snippet, which calculates the Intersection over Union (IoU) between two bounding boxes:

1 point

# Bounding box format: [x1, y1, x2, y2]

box1 = [50, 50, 150, 150]

box2 = [100, 100, 200, 200]

# Calculate the intersection coordinates

x1 = max(box1[0], box2[0])

y1 = max(box1[1], box2[1])

x2 = min(box1[2], box2[2])

y2 = min(box1[3], box2[3])

# Intersection area

intersection\_area = max(0, x2 - x1) \* max(0, y2 - y1)

# Areas of the two boxes

area\_box1 = (box1[2] - box1[0]) \* (box1[3] - box1[1])

area\_box2 = (box2[2] - box2[0]) \* (box2[3] - box2[1])

# Union area

union\_area = area\_box1 + area\_box2 - intersection\_area


# IoU calculation

iou = intersection\_area / union\_area if union\_area > 0 else 0

print(f"IoU: {iou:.2f}")

Given the bounding boxes [50,50,150,150] and [100,100,200,200], what is the calculated IoU?

☒ (a) 0.14



$$150 \times 200 = 30000$$

$$50 \times 100 = 5000$$

# Total loss

total\_loss = coord\_loss + object\_loss + no\_object\_loss + iou\_loss + class\_loss

print(f"Total Loss: {total\_loss:.4f}")

Which of the following loss components are included in the YOLO loss function illustrated in the code snippet?

☐ (a) Coordinate loss, Object loss, No-object loss, Classification loss.

☐ (b) Coordinate loss, Object loss, IoU loss, Classification loss.

☒ (c) Coordinate loss, IoU loss, Object loss, No-object loss, Classification loss.

☐ (d) IoU loss, Object loss, Classification loss, Regularization loss.

11) Which of the following networks can be used as feature extractors in Faster R-CNN?

1 point

☒ (a) VGG16, ResNet50, MobileNet

☐ (b) YOLO, SSD, RetinaNet

☐ (c) AlexNet, LeNet, GAN

☐ (d) RPN, FPN, DenseNet

No, the answer is incorrect.

Score: 0

Accepted Answers:

(a) VGG16, ResNet50, MobileNet

12) Which of the following statements correctly differentiates Fast R-CNN from Slow R-CNN?

1 point

☐ (a) Fast R-CNN uses external region proposals, while Slow R-CNN generates region proposals directly during training.

☐ (b) Slow R-CNN uses ROI Pooling for fixed-size feature extraction, while Fast R-CNN processes entire images in a single pass.

☒ (c) Fast R-CNN performs feature extraction on the entire image once, while Slow R-CNN extracts features separately for each region proposal.

☐ (d) Slow R-CNN integrates region proposal generation into the network, while Fast R-CNN relies on Selective Search for proposals.

13) YOLO uses a single neural network to predict bounding boxes and class probabilities directly from the entire image in one pass. (True/False)

true

14) In Faster R-CNN, the Region Proposal Network (RPN) doesn't share convolutional layers with the backbone network to generate region proposals, reducing computation time compared to external proposal methods. (True/False)

false